Fake or Real news



Done by: Lujain Yousef

Columns:

• title:

This column contains the titles of news articles or headlines. It represents the main headline or title of the new

title of the news story.

text:

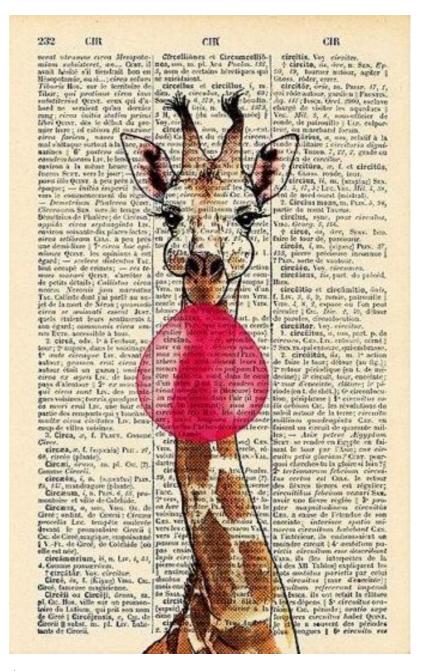
This column contains the main body of the news article. It includes the full text of the news story, providing more context and details beyond the headline.

· label:

This column indicates whether the news article is classified as "FAKE" or "REAL". It's likely a binary classification where "FAKE" denotes articles that are considered false or misleading, while "REAL" denotes articles that are considered accurate.

Text Vectorization and Classification:

 I performed text vectorization using both CountVectorizer and TF-IDF to convert the text data into numerical features suitable for machine learning models. After that, I applied Multinomial Naive Bayes classifier to classify the news articles into "FAKE" or "REAL" categories.



Data

```
import pandas as pd
df = pd.read_csv('fake_or_real_news.csv',index_col='Unnamed: 0')
df
```

Out[7]:		title	text	label
	8476	You Can Smell Hillary's Fear	Daniel Greenfield, a Shillman Journalism Fello	FAKE
	10294	Watch The Exact Moment Paul Ryan Committed Pol	Google Pinterest Digg Linkedin Reddit Stumbleu	FAKE
	3608	Kerry to go to Paris in gesture of sympathy	U.S. Secretary of State John F. Kerry said Mon	REAL
	10142	Bernie supporters on Twitter erupt in anger ag	— Kaydee King (@KaydeeKing) November 9, 2016 T	FAKE
	875	The Battle of New York: Why This Primary Matters	It's primary day in New York and front-runners	REAL
	•••			
	4490	State Department says it can't find emails fro	The State Department told the Republican Natio	REAL
	8062	The 'P' in PBS Should Stand for 'Plutocratic'	The 'P' in PBS Should Stand for 'Plutocratic'	FAKE
	8622	Anti-Trump Protesters Are Tools of the Oligarc	Anti-Trump Protesters Are Tools of the Oligar	FAKE
	4021	In Ethiopia, Obama seeks progress on peace, se	ADDIS ABABA, Ethiopia —President Obama convene	REAL
	4330	Jeb Bush Is Suddenly Attacking Trump. Here's W	Jeb Bush Is Suddenly Attacking Trump. Here's W	REAL

6335 rows × 3 columns

```
In [2]: df.columns
Out[2]: Index(['title', 'text', 'label'], dtype='object')
```

Named entity recognation

```
In [8]: from nltk import word_tokenize, pos_tag, ne_chunk
    from nltk.chunk import conlltags2tree, tree2conlltags
# Function to perform named entity recognition on text
def extract_entities(text):
    tokens = word_tokenize(text)
    pos_tags = pos_tag(tokens)
    tree = ne_chunk(pos_tags)
    iob_tags = tree2conlltags(tree)
    entities = [(token, iob) for token, _, iob in iob_tags if iob != 'O']
    return entities
```

```
In [9]: # Apply the function to the 'text_column' and create a new column for extracted endf['extracted_entities'] = df['text'].apply(extract_entities)
```

In [6]: display(df[['text','extracted_entities']])

	text	extracted_entities
8476	Daniel Greenfield, a Shillman Journalism Fello	[(Daniel, B-PERSON), (Greenfield, B-ORGANIZATI
10294	Google Pinterest Digg Linkedin Reddit Stumbleu	[(Google, B-PERSON), (Pinterest, B-PERSON), (D
3608	U.S. Secretary of State John F. Kerry said Mon	[(U.S., B-GPE), (State, B-ORGANIZATION), (John
10142	— Kaydee King (@KaydeeKing) November 9, 2016 T	[(Kaydee, B-PERSON), (King, I-PERSON), (Dem, B
875	It's primary day in New York and front-runners	[(New, B-GPE), (York, I-GPE), (Hillary, B-PERS
•••		
4490	The State Department told the Republican Natio	[(State, B-ORGANIZATION), (Department, I-ORGAN
8062	The 'P' in PBS Should Stand for 'Plutocratic'	[(PBS, B-ORGANIZATION), (Should, I-ORGANIZATIO
8622	Anti-Trump Protesters Are Tools of the Oligar	[(Oligarchy, B-ORGANIZATION), (Arthur, B-PERSO
4021	ADDIS ABABA, Ethiopia —President Obama convene	[(ADDIS, B-ORGANIZATION), (ABABA, B-GPE), (Eth
4330	Jeb Bush Is Suddenly Attacking Trump. Here's W	[(Jeb, B-PERSON), (Bush, B-ORGANIZATION), (Jeb

6335 rows × 2 columns

Words Frequency & dispersion plot

```
In [24]:
         import nltk
         from nltk.tokenize import word_tokenize
         from nltk.probability import FreqDist
         import string
         import matplotlib.pyplot as plt
         # Combine all text from the 'text' column into a single string
         all_text = ' '.join(df['text'])
         # Tokenize the combined text into words and remove punctuation
         tokens = word_tokenize(all_text)
         tokens = [word.lower() for word in tokens if word.isalnum()]
         # Remove stop words
         stop_words.update({'aa', 'aaa', '&vade', 'also', 'aaas'})
         filtered_tokens = [word for word in tokens if word not in stop_words]
         # Calculate the frequency distribution of words
         freq_dist = FreqDist(filtered_tokens)
```

```
# Select the most frequent words (e.g., top 10)
most_common_words = freq_dist.most_common(10)  # Adjust as needed

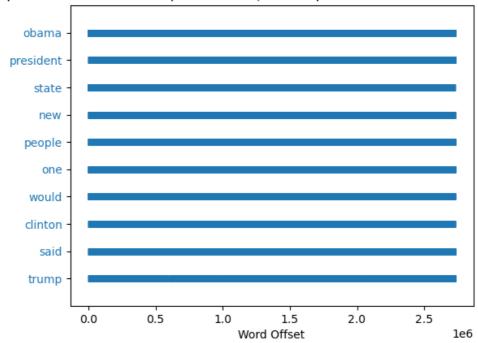
# Extract the most frequent words
words = [word for word, _ in most_common_words]

# Determine the positions of these words in the text
positions = [i for i, token in enumerate(filtered_tokens) if token in words]

# Create a dispersion plot
plt.figure(figsize=(10, 5))
nltk.draw.dispersion.dispersion_plot(filtered_tokens, words)
plt.title('Dispersion Plot of Most Frequent Words (with Stop Words and Punctuatio plt.show()
```

<Figure size 1000x500 with 0 Axes>

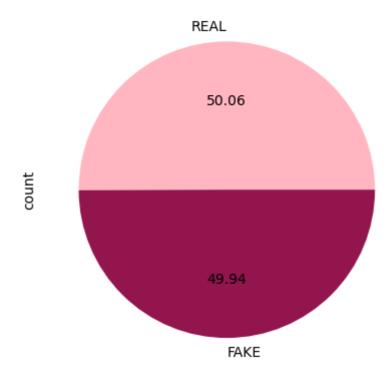
Dispersion Plot of Most Frequent Words (with Stop Words and Punctuation Removed)



Text Preprocessing

```
df.drop('extracted_entities',axis=1,inplace=True)
In [11]:
         from nltk.tokenize import RegexpTokenizer
In [21]:
         from nltk.stem import WordNetLemmatizer
         from nltk.corpus import stopwords
         def preprocess_text(text):
             tokenizer = RegexpTokenizer(r'\w+')
             tokens = tokenizer.tokenize(text.lower())
             # Lemmatize the words
             lemmatizer = WordNetLemmatizer()
             lemmatized_tokens = [lemmatizer.lemmatize(token) for token in tokens]
             # Remove stop words for English
             stop_words_en = set(stopwords.words('english'))
             stop_words_en.update({'aa', 'aaa', 'ยงade', 'also', 'aaas'})
             filtered_tokens = [token for token in tokens if token not in stop_words_en an
             # Remove stop words for Arabic
             stop_words_ar = set(stopwords.words('arabic'))
             filtered_tokens = [token for token in filtered_tokens if token not in stop_wo
             # Join the tokens back into a string
```

```
processed_text = ' '.join(filtered_tokens)
             return processed_text
In [22]: df['text'] = df['text'].apply( preprocess_text)
In [15]: df['text']
                  daniel greenfield shillman journalism fellow f...
         8476
Out[15]:
         10294
                  google pinterest digg linkedin reddit stumbleu...
         3608
                  u secretary state john f kerry said monday sto...
                  kaydee king kaydeeking november lesson tonight...
         10142
         875
                  primary day new york front runners hillary cli...
         4490
                  state department told republican national comm...
         8062
                  p pbs stand plutocratic pentagon posted oct wi...
         8622
                  anti trump protesters tools oligarchy reform a...
         4021
                  addis ababa ethiopia president obama convened ...
         4330
                  jeb bush suddenly attacking trump matters jeb ...
         Name: text, Length: 6335, dtype: object
         Splitting the Data
In [25]: X = df['text']
         y = df.label
In [26]: from sklearn.model_selection import train_test_split
         X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.30,random_state=
In [10]: y.value_counts()
         label
Out[10]:
         REAL
                 3171
         FAKE
                 3164
         Name: count, dtype: int64
 In [6]: import matplotlib.pyplot as plt
         # Assuming y is your pandas Series or DataFrame column
         y.value_counts().plot.pie(autopct='%.2f', colors=['#ffb6c1', "#93154d"])
         plt.show()
```



Count Vectorizer

```
from sklearn.feature_extraction.text import CountVectorizer
In [27]:
         count_vec = CountVectorizer()
In [28]: count_train = count_vec.fit_transform(X_train)
         count_test = count_vec.transform(X_test)
In [29]:
         # feature names
         count_vec.get_feature_names_out()[:10]
         array(['aab', 'aadmi', 'aahing', 'aaib', 'aalia', 'aam', 'aamaq', 'aamon',
Out[29]:
                'aap', 'aarhus'], dtype=object)
In [30]:
         import pandas as pd
         # Create DataFrame from CountVectorizer transformed data
         count_df = pd.DataFrame(count_train.A, columns=count_vec.get_feature_names_out())
         count df
```

Out[30]:		aab	aadmi	aahing	aaib	aalia	aam	aamaq	aamon	аар	aarhus	•••	الدولية	القادمون
	0	0	0	0	0	0	0	0	0	0	0		0	0
	1	0	0	0	0	0	0	0	0	0	0		0	0
	2	0	0	0	0	0	0	0	0	0	0		0	0
	3	0	0	0	0	0	0	0	0	0	0		0	0
	4	0	0	0	0	0	0	0	0	0	0		0	0
	•••													
	4429	0	0	0	0	0	0	0	0	0	0		0	0
	4430	0	0	0	0	0	0	0	0	0	0		0	0
	4431	0	0	0	0	0	0	0	0	0	0		0	0
	4432	0	0	0	0	0	0	0	0	0	0		0	0
	4433	0	0	0	0	0	0	0	0	0	0		0	0

4434 rows × 55393 columns

```
In [32]: count_df.columns

Out[32]: Index(['aab', 'aadmi', 'aahing', 'aaib', 'aalia', 'aam', 'aamaq', 'aamon', 'aap', 'aarhus',

...

'اللولية', 'اللولية', 'اللجنة', 'تحتاج', 'محاولات', 'والمرضى', 'والمرضى', 'والمرضى', 'والمرضى', 'dtype='object', length=48377)
```

Tfidf Vectorizer

```
In [31]: from sklearn.feature_extraction.text import TfidfVectorizer
    tfidf_vec = TfidfVectorizer()

In [32]: tfidf_train = tfidf_vec.fit_transform(X_train)
    tfidf_test = tfidf_vec.transform(X_test)

In [33]: tfidf_df = pd.DataFrame(tfidf_train.A,columns=tfidf_vec.get_feature_names_out())
    tfidf_df
```

Out[33]:		aab	aadmi	aahing	aaib	aalia	aam	aamaq	aamon	аар	aarhus	•••	الدولية	القادمون
	0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0		0.0	0.0
	1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0		0.0	0.0
	2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0		0.0	0.0
	3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0		0.0	0.0
	4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0		0.0	0.0
	•••		•••	•••				•••						
	4429	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0		0.0	0.0
	4430	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0		0.0	0.0
	4431	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0		0.0	0.0
	4432	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0		0.0	0.0
	4433	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0		0.0	0.0

4434 rows × 55393 columns

check count_df and tfidf_df if

```
In [34]: # Calculate the difference in columns: difference
    difference = set(count_df.columns) - set(tfidf_df.columns)
    print(difference)

# Check whether the DataFrames are equal
    print(count_df.equals(tfidf_df))
set()
False
```

Naive Bayes Model on count df

```
In [35]:
         from sklearn.pipeline import Pipeline
         from sklearn.model_selection import RandomizedSearchCV
         from sklearn.naive_bayes import MultinomialNB
         from scipy.stats import uniform, randint
         nb_classifier = MultinomialNB()
         # Define the pipeline with CountVectorizer and MultinomialNB
         pipeline = Pipeline([
              ('vect', count_vec),
              ('clf', nb_classifier )
         ])
         # Define the parameter distributions to sample from
         param dist = {
              'vect__min_df': randint(1, 3),
              'vect__max_df': uniform(0.5, 0.25), # Sampling from a uniform distribution b
              'clf__alpha': uniform(0.1, 0.9), # Sampling from a uniform distribution betw
         }
         # Perform randomized search using 5-fold cross-validation
         random_search = RandomizedSearchCV(pipeline, param_distributions=param_dist, n_it
```

```
random_search.fit(X_train, y_train)

# Print the best parameters found
print("Best parameters found:")
print(random_search.best_params_)

# Evaluate the model on the test set
accuracy = random_search.score(X_test, y_test)
print("Accuracy on test set:", accuracy)

Fitting 5 folds for each of 10 candidates, totalling 50 fits
Best parameters found:
{'clf_alpha': 0.3207004242339442, 'vect_max_df': 0.7004663059716135, 'vect_min_df': 2}
Accuracy on test set: 0.8979484481851657
```

Naive Bayes Model on tfidf df

```
In [36]: # Define the pipeline with TfidfVectorizer and MultinomialNB
         pipeline = Pipeline([
             ('vect', tfidf_vec),
             ('clf', nb_classifier)
         ])
         # Define the parameter distributions to sample from
         param_dist = {
             'vect__min_df': randint(1, 3),
             'vect__max_df': uniform(0.5, 0.25), # Sampling from a uniform distribution b
             'clf__alpha': uniform(0.1, 0.9), # Sampling from a uniform distribution betw
         }
         # Perform randomized search using 5-fold cross-validation
         random_search = RandomizedSearchCV(pipeline, param_distributions=param_dist, n_it
         random_search.fit(X_train, y_train)
         # Print the best parameters found
         print("Best parameters found:")
         print(random_search.best_params_)
         # Evaluate the model on the test set
         accuracy = random_search.score(X_test, y_test)
         print("Accuracy on test set:", accuracy)
         Fitting 5 folds for each of 10 candidates, totalling 50 fits
         Best parameters found:
         {'clf alpha': 0.23807810840304605, 'vect max df': 0.6592016762875144, 'vect mi
         n_df': 2}
         Accuracy on test set: 0.9021567596002105
```

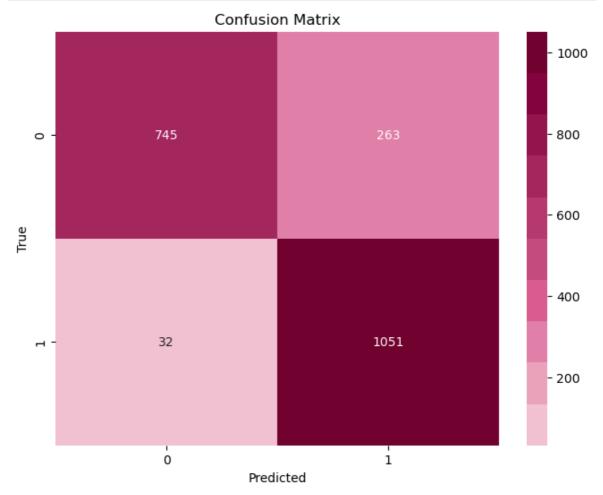
Confusion matrix

```
In [24]: y.unique()
Out[24]: array(['FAKE', 'REAL'], dtype=object)

In [55]: # Calculate the confusion matrix: cm
    cm = confusion_matrix(y_test, pred, labels=['FAKE', 'REAL'])

In [56]: # Plot the confusion matrix using a heatmap
    colors = ["#f1c1d1", "#eaa2bb", "#e080a8", "#d95c8f", "#c44b7e", "#b5386e", "#a32
```

```
# Plot the confusion matrix using a heatmap with the custom color palette
plt.figure(figsize=(8, 6))
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, cmap=colors, fmt='g', xticklabels=[0, 1], yticklabels
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```



The End